

CAPTCHA Bypass

Team 99

Nicholas Leung

Coden Mercurius

Pranavbhai Patel

Ravi Singh

Description

CAPTCHA, or Completely Automated Public Turing Test to Tell Computers and Humans Apart, is a challenge-response test that determines whether a user is authentic (human) or inauthentic (machine). They require users to authenticate themselves by retyping a character sequence prior to completing a request. This notebook implements a CAPTCHA bypass using deep learning. The team aims to investigate weaknesses and vulnerabilities of the CAPTCHA system.

```
In [185]: %%shell
jupyter nbconvert --to html /content/captcha.ipynb
```

```
[NbConvertApp] Converting notebook /content/captcha.ipynb to html
[NbConvertApp] Writing 695383 bytes to /content/captcha.html
```

Out[185]:

```
In [2]: # Imports
import torch
import torch.nn as nn
import os
from skimage import io
from torch.utils.data import Dataset, DataLoader
import torch.utils.data
import torchvision
from torchvision import datasets, transforms
import matplotlib.pyplot as plt
import numpy as np
import time
```

Part 1. Data Processing

The dataset for this model is generated using the following library: <https://github.com/lepture/captcha> (<https://github.com/lepture/captcha>) and automated by the script `dataset_generator.py` .

The character space began as purely numeric (0-9) but has since expanded to become alphanumeric (0-9, A-Z). Alphabetical characters are capitalized. Characters are uniformly distributed in terms of occurrence in the dataset.

The generated dataset `alphanumeric_dataset.zip` is available on the private team Google Drive because it is too large for the Github repository. Upload `alphanumeric_dataset.zip` into the Colab Files and unzip.

```
In [3]: # Unzip dataset
!unzip -qq /content/alphanumeric_dataset.zip -d /content/
```

```
In [5]: class CaptchaDataset(Dataset):
        """ Captcha Dataset """

        def __init__(self, directory):
            self.directory = directory
            self.captchas = os.listdir(directory)
            self.captchas.remove("metadata.txt")

            self.transform = transforms.Compose(
                [transforms.ToTensor(),
                 transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

            self.character_set = open(directory + "/metadata.txt", "r").readline().split(',')
            self.characters_to_identifier = {}

            for i in range(len(self.character_set)):
                self.characters_to_identifier.update({ self.character_set[i]: i })

        def __len__(self):
            # Assumes each file in the dataset directory represents a data sample
            return len(self.captchas)

        def __getitem__(self, index):
            sample_name = self.captchas[index]
            sample_captcha_values = list(sample_name[0:-4]) # Slice s.t. remove png file extension

            # Read the image and represent it as a tensor
            image = io.imread(self.directory + '/' + sample_name)
            image = self.transform(image)

            # Represent each character as an integer identifier
            label = []
            for char in sample_captcha_values:
                label.append(self.characters_to_identifier.get(char))

            return (image, torch.tensor(label))
```

```
In [6]: dataset_path = "/content/alphanumeric_dataset"

        # Instantiate dataset
        dataset = CaptchaDataset(dataset_path)
```

```
In [7]: def visualize_character_frequency(dataloader, title):
        character_frequency = {} # Contains frequency information
        character_set = dataset.character_set

        # Populate character_frequency
        for _, labels in dataloader:
            for label in labels:
                for char_identifier in label:
                    char = character_set[char_identifier.item()]
                    current_value = character_frequency.get(char, None)

                    if current_value is None:
                        character_frequency.update({ char : 0 })
                    else:
                        character_frequency.update({ char : current_value + 1 })

        x_values = range(len(character_set))
        y_values = []

        for char in character_set:
            count = character_frequency.get(char)
            y_values.append(count)

        plt.title(title)
        plt.plot(x_values, y_values)
        plt.xlabel("Characters")
        plt.ylabel("Count")
        plt.xticks(x_values, character_set)
        plt.show()
```

```
In [8]: def get_data_loaders(dataset, batch_size, total_size=None):

    if total_size is None:
        total_size = len(dataset)

    training_ratio = 0.7
    validation_ratio = 0.15
    # test_ratio implied

    train_length = int(total_size * training_ratio)
    validation_length = int((total_size - train_length) * (validation_ratio / (
1 - training_ratio)))
    test_length = total_size - train_length - validation_length
    fill = len(dataset) - total_size



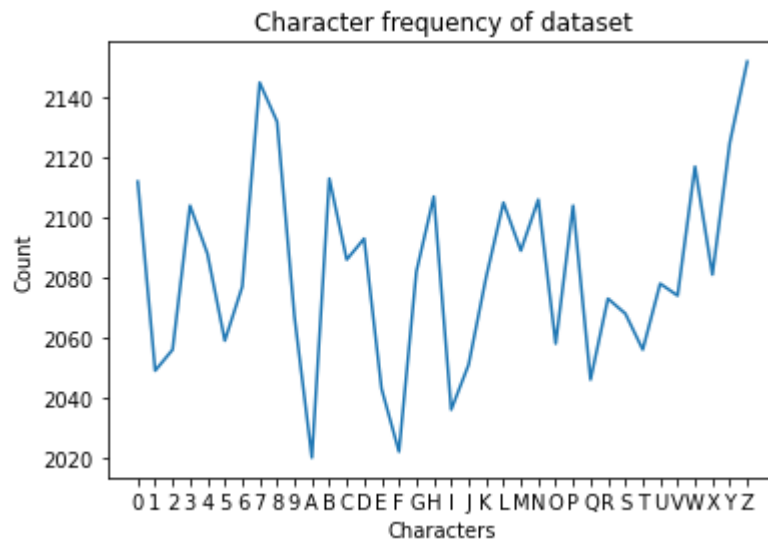
    train_set, valid_set, test_set, fill_set = torch.utils.data.random_split(dat
aset, [train_length, validation_length, test_length, fill], torch.Generator().
manual_seed(10))


    train_loader = torch.utils.data.DataLoader(train_set, batch_size=batch_size,
num_workers=1, drop_last=True, shuffle=True)
    valid_loader = torch.utils.data.DataLoader(valid_set, batch_size=batch_size,
num_workers=1, drop_last=True, shuffle=True)
    test_loader = torch.utils.data.DataLoader(test_set, batch_size=batch_size, n
um_workers=1, drop_last=True, shuffle=True)

    return train_loader, valid_loader, test_loader
```

```
In [9]: # Dataset visualization
train, valid, test = get_data_loaders(dataset, 100)

visualize_character_frequency(torch.utils.data.DataLoader(dataset, num_workers
=1), title="Character frequency of dataset")
visualize_character_frequency(train, title="Character frequency of training se
t")
visualize_character_frequency(valid, title="Character frequency of validation
set")
visualize_character_frequency(test, title="Character frequency of test set")
```





Part 2. Character Segmentation

Character segmentation must occur prior to character classification. This entails using `OpenCV.findContours()` to perform blob detection on a CAPTCHA image input. This will extract each individual character (5 total) for input to the model.

If characters are overlapping, the median width is assumed to separate the characters.

```
In [10]: import random
import cv2
import torchvision as tv
```



```
In [11]: def processimage(image, thresh):  
    #Format image type/ dimensions  
    image=image.permute(1,2,0)  
    image=image.numpy()  
    imageorig=image  
  
    #Modify image so contours/ borders can be easily found  
    #Greyscale  
    image = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)  
    #Binary Threshold  
    NA, image = cv2.threshold(image, thresh, 1, cv2.THRESH_BINARY)  
    #Erosion  
    kernel = np.ones((2,2),np.uint8)  
    image = cv2.dilate(image,kernel,iterations = 1)  
    #Vertical Blur and Resharpen  
    morpher = cv2.getStructuringElement(cv2.MORPH_RECT, (1,3))  
    image = cv2.morphologyEx(image, cv2.MORPH_CLOSE, morpher)  
    #Binary Threshold  
    thresh, image = cv2.threshold(image,thresh, 1, cv2.THRESH_BINARY)  
    #Expand Border  
    image=cv2.copyMakeBorder(image, 5, 5, 5, 5,cv2.BORDER_CONSTANT,value=1)  
    image = image.astype(np.uint8)  
  
    return imageorig,image
```

```

In [12]: def segmentimage(image,narrow):
    #Return list of borderlines in image
    contours, hierarchy = cv2.findContours(image, cv2.RETR_TREE, cv2.CHAIN_APPROX_SIMPLE)
    #imagecont=cv2.drawContours(image, contours, -1, (0, 0.5, 0), 1)

    #Creates boxes for every large object
    boxes=[]
    for contour in contours:
        [x,y,w,h]=cv2.boundingRect(contour)
        if(w>8 and w<120 and h>30):
            boxes.append([x,y,w,h])
    boxes.sort(key=lambda x: x[0])

    #Eliminates boxes that are contained within other boxes (subparts of a Letter)
    i=1
    while (i<len(boxes)):
        box=boxes[i]
        boxprev=boxes[i-1]
        if (box[0]>boxprev[0] and box[1]>boxprev[1] and (box[0]+box[2])<(boxprev[0]+boxprev[2]) and (box[1]+box[3])<(boxprev[1]+boxprev[3])):
            boxes.pop(i)
            i-=1
        i+=1

    #If boxes are too wide they may contain multiple boxes
    #They are split vertically into 2 or 3 subboxes (even width splits)
    i=0
    while (i<len(boxes)):
        box=boxes[i]
        if (box[2]>box[3]*(0.95-narrow)):
            x,y,w,h=boxes.pop(i)
            boxes.insert(i,[x+int((2*w)/3),y,int(w/3),h])
            boxes.insert(i,[x+int((w)/3),y,int(w/3),h])
            boxes.insert(i,[x,y,int(w/3),h])
        elif (box[2]>box[3]*(0.6-narrow)):
            x,y,w,h=boxes.pop(i)
            boxes.insert(i,[x+int(w/2),y,int(w/2),h])
            boxes.insert(i,[x,y,int(w/2),h])
        if (i>=len(boxes)-1):
            break
        i+=1

    return boxes

```

```

In [14]: def getcharacterimages(images, showsegments=False, filterBadSegmentation=False
, imgsize=80):
    characters=[]
    for i in range (0,len(images)):
        imageraw = images[i]

        # CAPTCHA image pre-processed, custom function called
        # Rectangle Borders of each character obtained, custom function called
        for i in range(0,7):
            thresh, narrow = 0.6,0
            if (i>=1):
                narrow=0.1
                thresh=0.6+(i-1)*0.1
            imageorig, image=processimage(imageraw,thresh)
            imageboxes = np.copy(image)
            boxes=segmentimage(image,narrow)
            if (len(boxes)>=5):
                break

        # Filter bad segmentation cases
        if filterBadSegmentation and len(boxes) < 5:
            continue

        # Individual Characters images are cut out from CAPTCHA image
        charactersset=[]
        for i in range(0,5):
            # If insufficient letters obtainable, add an empty image
            if (i<len(boxes)):
                box=boxes[i]
            else:
                box=[0,0,1,1]

            [x,y,w,h]=box
            char=image[y:y+h,x:x+w]
            height=char.shape[0]
            width=char.shape[1]
            # cv2.copyMakeBorder(soruce, top, bottom, Left, right, borderType, valu
e)
            if (width>height):
                char=cv2.copyMakeBorder(char, int ((width-height)/2),int ((width-heigh
t)/2), 0, 0,cv2.BORDER_CONSTANT,value=1)
            if (height>width):
                char=cv2.copyMakeBorder(char,0,0, int ((height-width)/2), int ((height
-width)/2),cv2.BORDER_CONSTANT,value=1)
            char = cv2.resize(char, dsize=(imgsize, imgsize), interpolation=cv2.INTER_CUBIC)
            char=torch.Tensor(char)
            charactersset.append(char)
            cv2.rectangle(imageboxes,(x,y),(x+w,y+h),0,1)
            charactersset=torch.stack(charactersset)
            characters.append(charactersset)

        if (showsegments==True):
            plt.imshow(imageorig)
            plt.show()
            plt.imshow(imageboxes, cmap='gray', vmin = 0, vmax = 1)

```

```
plt.show()
for i in range(0,5):
    plt.subplot(1,5,i+1)
    plt.imshow(charactersset[i], cmap='gray', vmin = 0, vmax = 1)
plt.show()

return torch.stack(characters)
```

```
In [15]: """
How To Use - getcharacterimages(images, showsegments=False)

Input = tensor(batchsize,numchannels,height,width) (see below)
Output = tensor(batchsize, numcharacters = 5, height = 80, width = 80 )

Set `showsegments` to `True` to visualize segmentation
"""

train, valid, test = get_data_loaders(dataset, 100)

images_processed = 0
successful_segmentation = 0

for images, labels in train:
    characters = getcharacterimages(images, False, True)
    images_processed = images_processed + len(labels)
    successful_segmentation = successful_segmentation + len(characters)

print(f"Successful segmentation accuracy: { successful_segmentation / images_p
rocessed }")
```

Successful segmentation accuracy: 0.9701923076923077

Part 3. Base Model

The base model is a non-deep learning method. The base model leverages the same character segmentation module and an SVM architecture is used for character classification.

The base model is to serve as a baseline of comparison for the primary model.

```
In [16]: from sklearn import svm
import numpy as np
```

```
In [17]: class BaseModel:

    def __init__(self):
        self.classifier = svm.SVC()

    def fit_classifier(self, dataloader):

        # Preprocessing to make our PyTorch data in acceptable format

        input_acc = []
        labels_acc = []

        for images, labels in dataloader:

            segmented_captchas = getcharacterimages(images, imgsize=28)

            # Iterate over each captcha
            for i in range(len(segmented_captchas)):
                captcha = segmented_captchas[i]

                # Iterate over each character
                for j in range(len(captcha)):
                    input_acc.append(captcha[j].detach().numpy().reshape(-1))
                    labels_acc.append(labels[i][j].detach().numpy())

            input_acc = np.array(input_acc)
            labels_acc = np.array(labels_acc)

            # Train character classification

            self.classifier.fit(input_acc, labels_acc)

    def predict(self, images):
        segmented_captchas = getcharacterimages(images, imgsize=28)

        output = []
        for captcha in segmented_captchas:

            out_captcha = []

            for character in captcha:
                numpy_char = character.detach().numpy().reshape((1, -1)) # Reshape to acceptable input for SVM predict()
                out_char = self.classifier.predict(numpy_char)
                out_captcha.append(out_char.item())

            output.append(out_captcha)

        return torch.tensor(output)
```

```
In [18]: model = BaseModel()
train_small, valid_small, test_small = get_data_loaders(dataset, 100, 3000)

model.fit_classifier(train_small)
```

```
In [21]: def evaluate_base_model(dataloader):

    total_character_guesses = 0
    total_captcha_guesses = 0

    incorrect_character_guesses = 0
    incorrect_captcha_guesses = 0

    failed_guess_frequency = {}

    for images, labels in dataloader:
        out = model.predict(images)

        # Iterate through each sample captcha in batch
        for i in range(len(labels)):
            bad_guess = False

            # Iterate through each character of captcha
            for j in range(len(labels[i])):

                total_character_guesses = total_character_guesses + 1
                guess = out[i][j]
                expected = labels[i][j]

                if (guess != expected):
                    incorrect_character_guesses = incorrect_character_guesses + 1

                    # Track per character bad guesses
                    current_failed_guess_count = failed_guess_frequency.get(dataset.character_set[guess], 0)
                    failed_guess_frequency.update({ dataset.character_set[guess]: current_failed_guess_count + 1 })

                    bad_guess = True

            if bad_guess:
                incorrect_captcha_guesses = incorrect_captcha_guesses + 1

        total_captcha_guesses = total_captcha_guesses + 1

    # Overall accuracy information

    character_guess_accuracy = (total_character_guesses - incorrect_character_guesses) / total_character_guesses
    captcha_guess_accuracy = (total_captcha_guesses - incorrect_captcha_guesses) / total_captcha_guesses

    print(f"Character Accuracy: {character_guess_accuracy}")
    print(f"Captcha Accuracy: {captcha_guess_accuracy}")

    # Plot incorrect character guess frequency

    bad_guess_character_set = failed_guess_frequency.keys()

    x_values = range(len(bad_guess_character_set))
    y_values = []
```

```

for char in bad_guess_character_set:
    count = failed_guess_frequency.get(char)
    y_values.append(count)


plt.title("Bad guess character frequency")
plt.plot(x_values, y_values)
plt.xlabel("Characters")
plt.ylabel("Count")
plt.xticks(x_values, bad_guess_character_set)
plt.show()

```

In [22]: evaluate_base_model(valid)

Character Accuracy: 0.8199090909090909

Captcha Accuracy: 0.5095454545454545



Part 4: Primary Model

The primary model is a standard CNN with two convolutional layers and two fully-connected layers. Each convolutional layer is coupled with a max pooling layer (stride 2).

```

In [27]: import torch.nn.functional as F
import torch.optim as optim
import pandas as pd

```

```
In [28]: # Sample CNN based on Lab 2 model
class CaptchaCNN(nn.Module):
    def __init__(self):
        super(CaptchaCNN, self).__init__()
        self.name = "CaptchaCNN"

        # 5 chars, 5,5 we decide
        self.conv1 = nn.Conv2d(1, 5, 5)
        # w,h =
        self.pool1 = nn.MaxPool2d(2, 2)
        # w,h =

        # 5 matches the second 5 from the first CNN layer
        self.conv2 = nn.Conv2d(5, 10, 5)
        # w,h =
        self.pool2 = nn.MaxPool2d(2,2)
        # w,h =

        self.fc1 = nn.Linear(160, 32)
        self.fc2 = nn.Linear(32, 36)

    def forward(self, img):
        x = getcharacterimages(img, False, imgsize=28)
        x = x.reshape(-1, 1, 28, 28)
        x = self.pool1(F.relu(self.conv1(x)))
        x = self.pool2(F.relu(self.conv2(x)))
        x = x.view(-1, 160)
        return self.fc2(F.relu(self.fc1(x)))
```

```
In [29]: def plot(title, xlabel, ylabel, data1, data1_label, data2, data2_label, epochs
):
    plt.title(title)
    plt.plot(epochs, data1, label=data1_label)
    if data2 is not None:
        plt.plot(epochs, data2, label=data2_label)
    plt.xlabel(xlabel)
    plt.ylabel(ylabel)
    plt.legend(loc='best')
    plt.show()
```



```
In [30]: def get_accuracy(model, data_loader):
    total = 0
    correct = 0
    char_correct = 0
    captcha_length = 5
    for imgs, labels in data_loader:
        if torch.cuda.is_available():
            imgs = imgs.cuda()
            labels = labels.cuda()
        for i in range(batch_size):
            output = model(imgs[i].unsqueeze(dim=0))
            pred = output.max(1, keepdim=True)[1] # get the index of the max Logit
            num_correct = 0
            for j in range(captcha_length):
                # print(labels[i][j])
                # print(pred[j])
                if labels[i][j] == pred[j].squeeze(0):
                    num_correct += 1
                    char_correct += 1
                # print(num_correct)
            if num_correct == 5:
                correct += 1
            total += 1
    return correct / total, char_correct / (5 * total)
```

```
In [31]: def get_model_name(name, epoch, learning_rate=1e-4):
    """
    Generate a name for the model consisting of all the hyperparameter values
    """
    path="model_{0}_lr{1}_epoch{2}".format(
        name, learning_rate, epoch)
    return path
```

```

In [106]: def train_cnn(model, x, y, batch_size=128, num_epochs=20, learning_rate=0.001
):
    torch.manual_seed(360)

    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=learning_rate)

    train_loader = x
    valid_loader = y

    iters = []
    losses = []
    train_acc = []
    valid_acc = []

    epoch = 0

    for epoch in range(num_epochs):
        for imgs, labels in train_loader:
            labels = labels.reshape(-1)
            if use_cuda and torch.cuda.is_available():
                imgs = imgs.cuda()
                labels = labels.cuda()
            out = model(imgs) # forward pass
            loss = criterion(out, labels) # compute the total loss
            loss.backward() # backward pass (compute parameter updates)
            optimizer.step() # make the updates for each parameter
            optimizer.zero_grad() # a clean up step for PyTorch

            # Save the current model (checkpoint) to a file
            model_path = get_model_name(model.name, epoch, learning_rate=learning_
rate)
            torch.save(model.state_dict(), model_path)

            # save the current training information
            iters.append(epoch)
            losses.append(float(loss)) # compute *average* loss
            captcha_acc, char_acc = get_accuracy(model, train_loader)
            train_acc.append(captcha_acc) # compute training accuracy
            if y != None:
                valid_acc.append(
                    get_accuracy(model, valid_loader)[0]
                ) # compute validation accuracy
            if y != None:
                print(
                    (
                        "Epoch {}: Character accuracy: {}, Training accuracy: {},
" + "Validation accuracy: {}"
                    ).format(epoch + 1, char_acc, train_acc[epoch], valid_acc[epoch])
                )
            else:
                print(
                    ("Epoch {}: Character accuracy: {}, Training accuracy: {}".fo
rmat(epoch + 1, char_acc, train_acc[epoch])
                )

```

```
        epoch += 1

plt.title("Loss")
plt.plot(iters, losses, label="Train")
plt.xlabel("Iterations")
plt.ylabel("Loss")
plt.show()

plt.title("Accuracy")
plt.plot(iters, train_acc, label="Train")
if y != None:
    plt.plot(iters, valid_acc, label="Validation")
plt.xlabel("Iterations")
plt.ylabel("Accuracy")
plt.legend(loc="best")
plt.show()

if y != None:
    return losses, train_acc, valid_acc
return losses, train_acc
```


```
In [107]: model = CaptchaCNN()  
          train, valid, test = get_data_loaders(dataset, 32)  
          losses, train_acc, valid_acc = train_cnn(model, train, valid, 32, 30, 0.001)
```

Epoch 1: Character accuracy: 0.685766006097561, Training accuracy: 0.20807926
829268292, Validation accuracy: 0.2080357142857143
Epoch 2: Character accuracy: 0.7602705792682927, Training accuracy: 0.3469893
292682927, Validation accuracy: 0.3513392857142857
Epoch 3: Character accuracy: 0.7902248475609757, Training accuracy: 0.4193025
9146341464, Validation accuracy: 0.4316964285714286
Epoch 4: Character accuracy: 0.8091844512195122, Training accuracy: 0.4643673
780487805, Validation accuracy: 0.4660714285714286
Epoch 5: Character accuracy: 0.8171684451219512, Training accuracy: 0.4868521
341463415, Validation accuracy: 0.48526785714285714
Epoch 6: Character accuracy: 0.8233612804878049, Training accuracy: 0.5041920
731707317, Validation accuracy: 0.5022321428571429
Epoch 7: Character accuracy: 0.8323170731707317, Training accuracy: 0.5270579
268292683, Validation accuracy: 0.5178571428571429
Epoch 8: Character accuracy: 0.8409870426829268, Training accuracy: 0.5521150
914634146, Validation accuracy: 0.5410714285714285
Epoch 9: Character accuracy: 0.8378239329268292, Training accuracy: 0.5402057
926829268, Validation accuracy: 0.5241071428571429
Epoch 10: Character accuracy: 0.8426067073170732, Training accuracy: 0.556592
9878048781, Validation accuracy: 0.5339285714285714
Epoch 11: Character accuracy: 0.8457507621951219, Training accuracy: 0.565358
231707317, Validation accuracy: 0.5513392857142857
Epoch 12: Character accuracy: 0.8515434451219512, Training accuracy: 0.584889
481707317, Validation accuracy: 0.5683035714285715
Epoch 13: Character accuracy: 0.8530678353658536, Training accuracy: 0.587652
4390243902, Validation accuracy: 0.5736607142857143
Epoch 14: Character accuracy: 0.8548208841463415, Training accuracy: 0.594035
8231707317, Validation accuracy: 0.5727678571428572
Epoch 15: Character accuracy: 0.8534679878048781, Training accuracy: 0.588700
4573170732, Validation accuracy: 0.565625
Epoch 16: Character accuracy: 0.8560403963414634, Training accuracy: 0.594798
018292683, Validation accuracy: 0.5834821428571428
Epoch 17: Character accuracy: 0.8531440548780488, Training accuracy: 0.584032
0121951219, Validation accuracy: 0.5589285714285714
Epoch 18: Character accuracy: 0.8606897865853659, Training accuracy: 0.611756
8597560976, Validation accuracy: 0.5910714285714286
Epoch 19: Character accuracy: 0.8612804878048781, Training accuracy: 0.611185
2134146342, Validation accuracy: 0.596875
Epoch 20: Character accuracy: 0.8640625, Training accuracy: 0.62214176829268
3, Validation accuracy: 0.5946428571428571
Epoch 21: Character accuracy: 0.863795731707317, Training accuracy: 0.6199504
573170732, Validation accuracy: 0.6008928571428571
Epoch 22: Character accuracy: 0.8588414634146342, Training accuracy: 0.601181
4024390244, Validation accuracy: 0.5808035714285714
Epoch 23: Character accuracy: 0.8660251524390243, Training accuracy: 0.629382
6219512195, Validation accuracy: 0.5924107142857142
Epoch 24: Character accuracy: 0.8645198170731707, Training accuracy: 0.621189
0243902439, Validation accuracy: 0.5959821428571429
Epoch 25: Character accuracy: 0.867701981707317, Training accuracy: 0.6315739
329268293, Validation accuracy: 0.5991071428571428
Epoch 26: Character accuracy: 0.8661204268292683, Training accuracy: 0.625285
8231707317, Validation accuracy: 0.5870535714285714
Epoch 27: Character accuracy: 0.8681021341463414, Training accuracy: 0.631955
0304878049, Validation accuracy: 0.6013392857142857
Epoch 28: Character accuracy: 0.8666920731707317, Training accuracy: 0.629192
0731707317, Validation accuracy: 0.5924107142857142
Epoch 29: Character accuracy: 0.8721989329268293, Training accuracy: 0.645007

6219512195, Validation accuracy: 0.609375

Epoch 30: Character accuracy: 0.8690167682926829, Training accuracy: 0.635861

2804878049, Validation accuracy: 0.5977678571428572



Part 5: AlexNet Transfer Learning

The pretrained AlexNet model was imported as a backup to the primary model if the CNN architecture yielded subpar results. Similar to the primary model, the AlexNet model was coupled with the character segmentation module.

```
In [73]: import torchvision.models
alexnet = torchvision.models.alexnet(pretrained=True)

AlexNet_train, AlexNet_valid = [], []

train, valid, test = get_data_loaders(dataset, 100)

for images, labels in train:
    AlexNet_train.append((getcharacterimages(images, False, imgsize=80), labels))

for images, labels in valid:
    AlexNet_valid.append((getcharacterimages(images, False, imgsize=80), labels))
```

```
In [74]: imgs_train, labels_train = [], []
         imgs_valid, labels_valid = [], []

for img, label in AlexNet_train:
    img_grey = img.reshape(-1, 1, 80, 80)
    label = label.reshape(-1)
    img_color = img_grey.repeat(1,3,1,1)
    features = torch.from_numpy(alexnet.features(img_color).detach().numpy())
    imgs_train.append(features)
    labels_train.append(label)

for img, label in AlexNet_valid:
    img = img.reshape(-1, 1, 80, 80)
    label = label.reshape(-1)
    img_color = img.repeat(1,3,1,1)
    features = torch.from_numpy(alexnet.features(img_color).detach().numpy())
    imgs_valid.append(features)
    labels_valid.append(label)

AlexNet_train = list(zip(imgs_train, labels_train))
AlexNet_valid = list(zip(imgs_valid, labels_valid))
```

```

In [60]: def get_accuracy_alexnet(model, data_loader):
    total = 0
    correct = 0
    char_correct = 0
    captcha_length = 5
    for imgs, labels in data_loader:
        if torch.cuda.is_available():
            imgs = imgs.cuda()
            labels = labels.cuda()
        output = model(imgs)
        pred = output.max(1, keepdim=True)[1].squeeze(1) # get the index of the
max logit
        for i in range(0, len(output), captcha_length):
            num_correct = 0
            for j in range(captcha_length):
                if labels[i+j] == pred[i+j]:
                    num_correct += 1
                    char_correct += 1
            if num_correct == 5:
                correct += 1
            total += 1
    return correct / total, char_correct / (5 * total)

def train_alexnet(model, x, y, batch_size=128, num_epochs=20, learning_rate=0.
001):
    torch.manual_seed(360)

    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=learning_rate)

    train_loader = x
    valid_loader = y

    iters = []
    losses = []
    train_acc = []
    valid_acc = []

    epoch = 0

    for epoch in range(num_epochs):
        for img, label in train_loader:
            if use_cuda and torch.cuda.is_available():
                img = img.cuda()
                label = label.cuda()
            out = model(img) # forward pass
            loss = criterion(out, label) # compute the total loss
            loss.backward() # backward pass (compute parameter updates)
            optimizer.step() # make the updates for each parameter
            optimizer.zero_grad() # a clean up step for PyTorch

            # Save the current model (checkpoint) to a file
            model_path = get_model_name(model.name, epoch, learning_rate=learning_
rate)
            torch.save(model.state_dict(), model_path)

```



```

# save the current training information
iters.append(epoch)
losses.append(float(loss)) # compute *average* loss
captcha_acc, char_acc = get_accuracy_alexnet(model, train_loader)
train_acc.append(captcha_acc) # compute training accuracy
if y != None:
    captcha_acc, char_acc = get_accuracy_alexnet(model, valid_loader)
# compute validation accuracy
valid_acc.append(captcha_acc)
if y != None:
    print(
        (
            "Epoch {}: Character accuracy: {}, Training accuracy: {},
" + "Validation accuracy: {}"
        ).format(epoch + 1, char_acc, train_acc[epoch], valid_acc[epoch])
    )
else:
    print(
        ("Epoch {}: Character accuracy: {}, Training accuracy: {}".format(
            epoch + 1, char_acc, train_acc[epoch]
        ))
    )
    epoch += 1

# model_path = get_model_name(model.name, batch_size, learning_rate, epoch)
# torch.save(model.state_dict(), model_path)
plt.title("Loss")
plt.plot(iters, losses, label="Train")
plt.xlabel("Iterations")
plt.ylabel("Loss")
plt.show()

plt.title("Accuracy")
plt.plot(iters, train_acc, label="Train")
if y != None:
    plt.plot(iters, valid_acc, label="Validation")
plt.xlabel("Iterations")
plt.ylabel("Accuracy")
plt.legend(loc="best")
plt.show()

if y != None:
    return losses, train_acc, valid_acc
return losses, train_acc

```

```
In [56]: class AlexNetANNClassifier(nn.Module):
    def __init__(self):
        super(AlexNetANNClassifier, self).__init__()
        self.name = "AlexNetANNClassifier"
        self.fc1 = nn.Linear(256, 128)
        self.fc2 = nn.Linear(128, 64)
        self.fc3 = nn.Linear(64, 36)

    def forward(self, img):
        x = img.view(-1, 256)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        x = F.relu(x)
        x = self.fc3(x)
        return x
```


```
In [75]: model = AlexNetANNCNNClassifier()
         if torch.cuda.is_available():
             model.cuda()

         losses, train_acc, valid_acc = train_alexnet(model, AlexNet_train, AlexNet_val
         id, 32, 30, 0.001)
```

Epoch 1: Character accuracy: 0.6286363636363637, Training accuracy: 0.14115384615384616, Validation accuracy: 0.1340909090909091
Epoch 2: Character accuracy: 0.7095454545454546, Training accuracy: 0.24548076923076922, Validation accuracy: 0.23772727272727273
Epoch 3: Character accuracy: 0.7362727272727273, Training accuracy: 0.2968269230769231, Validation accuracy: 0.29409090909090907
Epoch 4: Character accuracy: 0.7480909090909091, Training accuracy: 0.3325961538461538, Validation accuracy: 0.315
Epoch 5: Character accuracy: 0.7565454545454545, Training accuracy: 0.355, Validation accuracy: 0.33
Epoch 6: Character accuracy: 0.7648181818181818, Training accuracy: 0.3730769230769231, Validation accuracy: 0.35363636363636364
Epoch 7: Character accuracy: 0.7711818181818182, Training accuracy: 0.385, Validation accuracy: 0.36954545454545457
Epoch 8: Character accuracy: 0.7744545454545455, Training accuracy: 0.3978846153846154, Validation accuracy: 0.38181818181818183
Epoch 9: Character accuracy: 0.7774545454545455, Training accuracy: 0.40625, Validation accuracy: 0.3881818181818182
Epoch 10: Character accuracy: 0.7785454545454545, Training accuracy: 0.4124038461538462, Validation accuracy: 0.3936363636363636
Epoch 11: Character accuracy: 0.7813636363636364, Training accuracy: 0.4201923076923077, Validation accuracy: 0.3986363636363636
Epoch 12: Character accuracy: 0.7835454545454545, Training accuracy: 0.42567307692307693, Validation accuracy: 0.4
Epoch 13: Character accuracy: 0.7841818181818182, Training accuracy: 0.4322115384615385, Validation accuracy: 0.3990909090909091
Epoch 14: Character accuracy: 0.786, Training accuracy: 0.4355769230769231, Validation accuracy: 0.4036363636363636
Epoch 15: Character accuracy: 0.7872727272727272, Training accuracy: 0.4411538461538462, Validation accuracy: 0.4090909090909091
Epoch 16: Character accuracy: 0.7886363636363637, Training accuracy: 0.4489423076923077, Validation accuracy: 0.41409090909090907
Epoch 17: Character accuracy: 0.7895454545454546, Training accuracy: 0.4514423076923077, Validation accuracy: 0.41545454545454547
Epoch 18: Character accuracy: 0.7895454545454546, Training accuracy: 0.4558653846153846, Validation accuracy: 0.41545454545454547
Epoch 19: Character accuracy: 0.7894545454545454, Training accuracy: 0.45865384615384613, Validation accuracy: 0.415
Epoch 20: Character accuracy: 0.7883636363636364, Training accuracy: 0.4633653846153846, Validation accuracy: 0.4109090909090909
Epoch 21: Character accuracy: 0.7885454545454545, Training accuracy: 0.46826923076923077, Validation accuracy: 0.41363636363636364
Epoch 22: Character accuracy: 0.7896363636363637, Training accuracy: 0.47125, Validation accuracy: 0.41454545454545455
Epoch 23: Character accuracy: 0.7906363636363636, Training accuracy: 0.4794230769230769, Validation accuracy: 0.41818181818181815
Epoch 24: Character accuracy: 0.7922727272727272, Training accuracy: 0.485, Validation accuracy: 0.42
Epoch 25: Character accuracy: 0.7915454545454546, Training accuracy: 0.4901923076923077, Validation accuracy: 0.41954545454545455
Epoch 26: Character accuracy: 0.7916363636363636, Training accuracy: 0.4929807692307692, Validation accuracy: 0.41863636363636364
Epoch 27: Character accuracy: 0.7927272727272727, Training accuracy: 0.49673076923076925, Validation accuracy: 0.42363636363636364
Epoch 28: Character accuracy: 0.7930909090909091, Training accuracy: 0.500576923076923, Validation accuracy: 0.42363636363636364
Epoch 29: Character accuracy: 0.7932727272727272, Training accuracy: 0.505576

923076923, Validation accuracy: 0.4240909090909091

Epoch 30: Character accuracy: 0.7940909090909091, Training accuracy: 0.50875,
Validation accuracy: 0.425



Part 6: End-to-End System

The end-to-end system takes in a single CAPTCHA image and attempts to decode it. Currently, the end-to-end system performs character classification using the "best" CNN model from training as it yielded the highest accuracies without overfitting. The end-to-end system outputs its prediction of each character (5 total) and the input itself for comparison.

```
In [96]: def decodeCharacter (encodedValue):  
         if (encodedValue < 10):  
             return str(encodedValue)  
         else:  
             return chr(encodedValue + 55)
```

```
In [183]: def e2emodel (imgs):  
    model = CaptchaCNN()  
    model.load_state_dict(torch.load(get_model_name(model.name, epoch=20, learning_rate=0.001)))  
  
    # Prediction  
    out = model(imgs)  
    pred = out.max(1, keepdim=True)[1].squeeze(1).tolist()[:5]  
    CAPTCHA_prediction = list(map(decodeCharacter, pred))  
    print(f"Predicted output = {CAPTCHA_prediction}")  
  
    # Plot batch of images  
    plt.imshow(imgs[0][0])
```

```
In [184]: imgs, labels = next(iter(test))  
e2emodel(imgs)
```

Predicted output = ['L', 'V', 'E', 'Z', 'H']

